



Indian Institute of Technology Gandhinagar

---

*CS331 - Computer Networks*

# ASSIGNMENT 1

# REPORT

---

Name - Bhavya Parmar, Romit Mohane

Roll No. - 23110059, 23110279

Date of Submission - 15 Sept 2025

Link to [GitHub Repository](#)

---

# Table of Contents

---

Task 1 - DNS Resolver.....	1
Introduction & Setup.....	1
Procedure.....	1
Result.....	1
Task 2 - Traceroute Protocol Behavior.....	3
Introduction & Setup.....	3
Procedure.....	3
Questions.....	4

# Task 1 - DNS Resolver

---

## Introduction & Setup

The purpose of this task is to understand and implement packet parsing and processing logic in a programming language, specifically a custom DNS resolution support, and also give an idea of load balancing and socket programming.

## Procedure

After adding the last 3 digits of our roll numbers (279 + 059), we had to select the file '**8.pcap**'. We used the `struct` Python module to read the packed binary data of each packet and process the DNS packets from the `.pcap` file in our client.

The client, after appending the custom header to the top, sends the DNS query packets to the server's PORT using the `socket` module.

The server listens on this PORT for the incoming packets. On receiving a packet, it parses the header, and resolves the IP using the given rules. It then returns the resolved IP to the client.

The client logs these requests and responses in a `.txt` file.

## Result

Running the Client and Server at 4:14 AM (morning slot), we get the following:

Custom header value (HHMMSSID)	Domain name	Resolved IP address
04144000	github.com.	192.168.1.1
04144001	bing.com.	192.168.1.2
04144102	facebook.com.	192.168.1.3
04144103	amazon.com.	192.168.1.4
04144104	linkedin.com.	192.168.1.5
04144105	stackoverflow.com.	192.168.1.1

Running the Client and Server at 2:52 PM (afternoon slot), we get the following:

<b>Custom header value (HHMMSSID)</b>	<b>Domain name</b>	<b>Resolved IP address</b>
14525600	github.com.	192.168.1.6
14525601	bing.com.	192.168.1.7
14525702	facebook.com.	192.168.1.8
14525703	amazon.com.	192.168.1.9
14525804	linkedin.com.	192.168.1.10
14525905	stackoverflow.com.	192.168.1.6

Running the Client and Server at 11:22 PM (night slot), we get the following:

<b>Custom header value (HHMMSSID)</b>	<b>Domain name</b>	<b>Resolved IP address</b>
23222000	github.com.	192.168.1.11
23222001	bing.com.	192.168.1.12
23222102	facebook.com.	192.168.1.13
23222103	amazon.com.	192.168.1.14
23222104	linkedin.com.	192.168.1.15
23222105	stackoverflow.com.	192.168.1.11

## Task 2 - Traceroute Protocol Behavior

### Introduction & Setup

The purpose of this task is to understand how the traceroute utility works in different operating systems by capturing the network traffic during the executions using Wireshark.

We chose to work on Windows and Linux. And we decided to trace the route to [www.dev.to](http://www.dev.to).

### Procedure

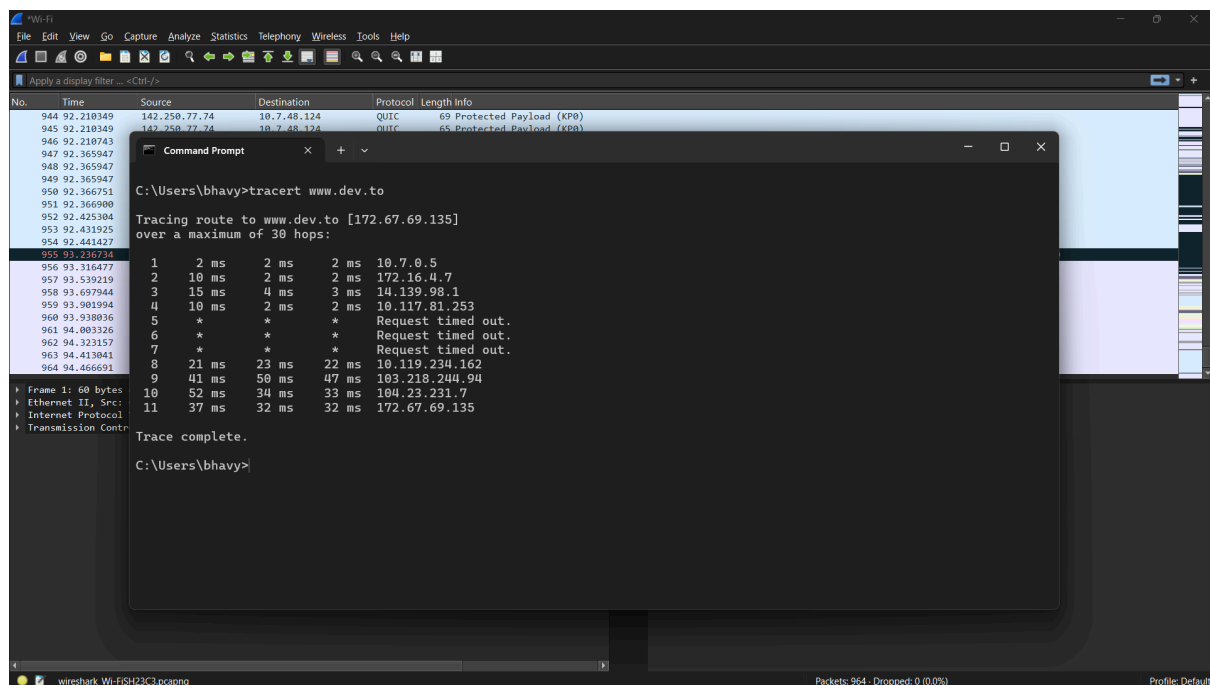
#### Windows

First, we opened up Wireshark and selected WiFi to capture, following which we ran the command:

```
tracert www.dev.to
```

This started to trace the route to the given destination over a maximum of 30 hops.

Upon completion, I stopped the packet capture in Wireshark and saved the file in .pcapng format.



## Linux

In Linux, the procedure was similar to that of Windows... We opened Wireshark and then selected wlan0 (which is the WiFi interface) to capture. Following which we ran the command:

```
tracert www.dev.to
```

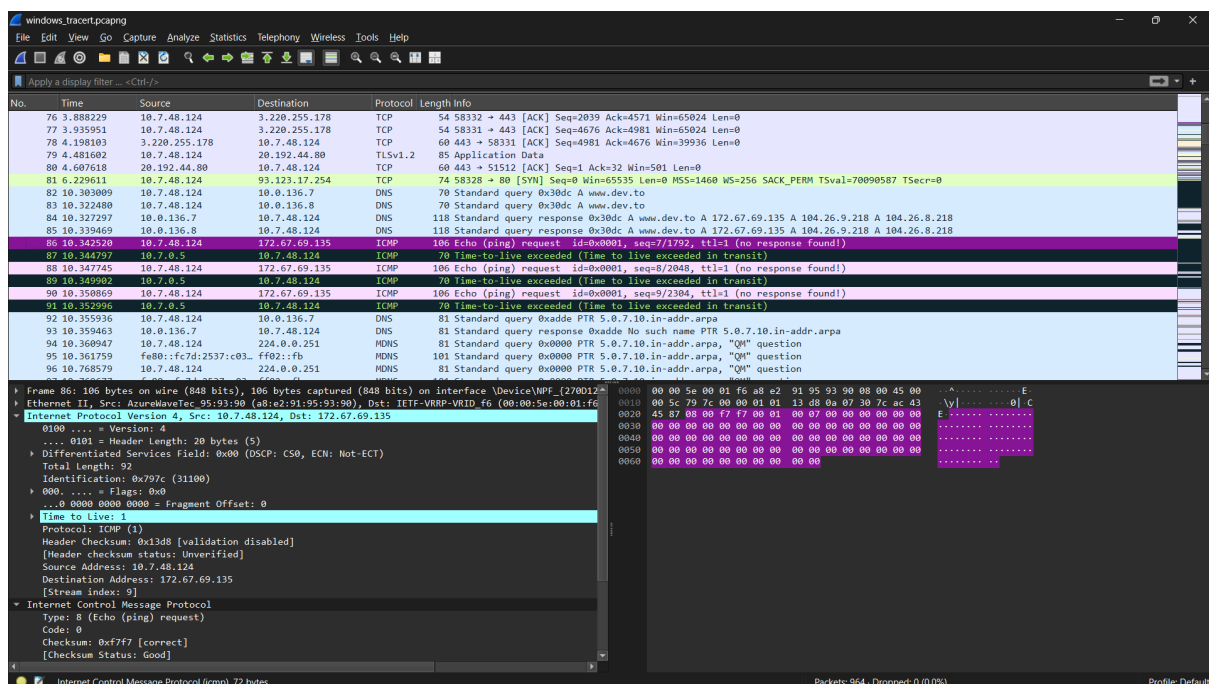
This started to trace the route to the given target over a maximum of 30 hops. And upon completion of this, I stopped the packet capture in Wireshark and saved the file in .pcapng format.

```
> traceroute www.dev.to
traceroute to www.dev.to (172.67.69.135), 30 hops max, 60 byte packets
 1  10.7.0.5 (10.7.0.5)  2.352 ms  2.180 ms  3.247 ms
 2  172.16.4.7 (172.16.4.7)  2.059 ms  2.000 ms  1.945 ms
 3  14.139.98.1 (14.139.98.1)  5.177 ms  5.046 ms  4.979 ms
 4  10.117.81.253 (10.117.81.253)  2.743 ms  2.690 ms  2.597 ms
 5  * * *
 6  * * *
 7  * * *
 8  10.119.234.162 (10.119.234.162)  24.501 ms  26.807 ms  26.757 ms
 9  103.218.244.94 (103.218.244.94)  37.680 ms  37.607 ms  38.709 ms
10  104.23.231.7 (104.23.231.7)  38.716 ms  104.23.231.5 (104.23.231.5)  34.810 ms  104.23.231.7 (104.23.231.7)  34.238 ms
11  172.67.69.135 (172.67.69.135)  36.295 ms  33.222 ms  36.856 ms
```

## Questions

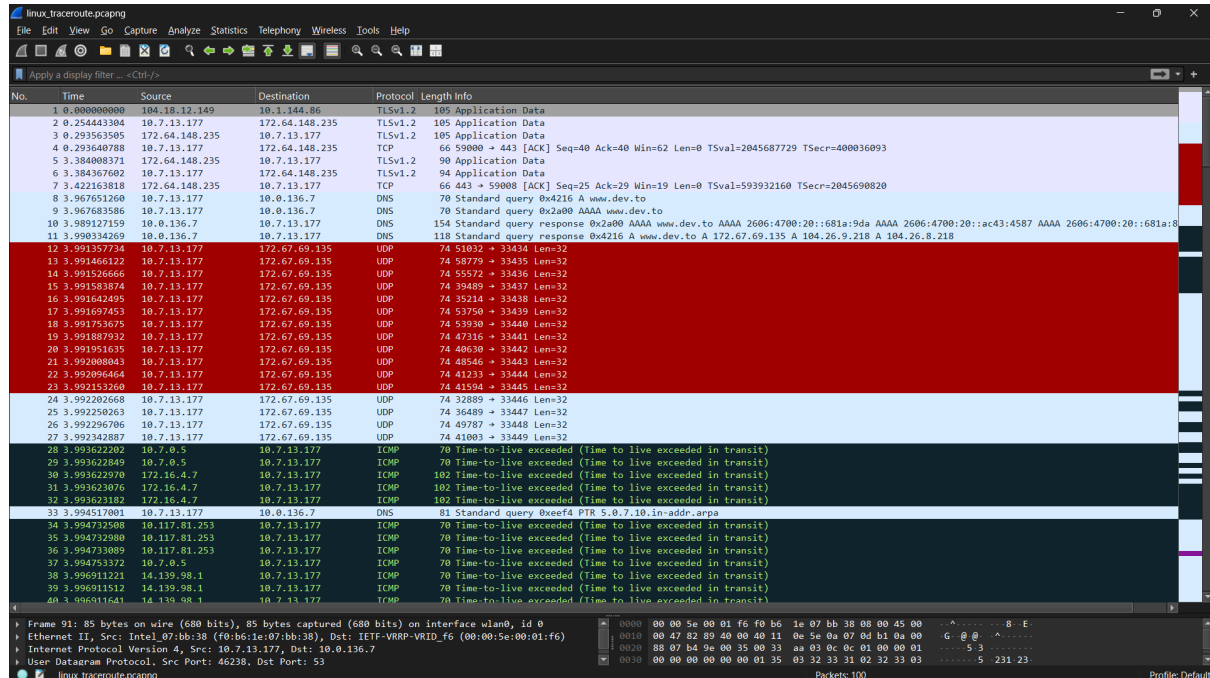
1. What protocol does Windows tracert use by default, and what protocol does Linux traceroute use by default?

Ans. Windows tracert uses the ICMP (Internet Control Message Protocol) by default, and sends echo request packets to the destination IP, probe-wise (3 messages in one probe), with the TTL (Time To Live) increasing at each probe. The intermediate routers respond with ICMP “Time Exceeded” messages.



The selected 86th packet is the first packet sent by the tracert operation via the ICMP protocol in the above image.

Whereas Linux traceroute uses UDP (User Datagram Protocol) by default, whilst also sending packets probe-wise (3 messages in one probe) to high-numbered ports (starting at 33434). Intermediate routers reply with ICMP "Time-to-live exceeded," and the destination replies with ICMP "Port Unreachable" (since no service listens on those high ports). The image showing the pcap file for Linux for the same is below:



**2. Some hops in your traceroute output may show \*\*\*. Provide at least two reasons why a router might not reply.**

Ans. In my tracert to www.dev.to, hops 5, 6, and 7 displayed \*\*\*, indicating that no response was received from those routers. This does not mean the route was broken, since subsequent hops (8 onwards) continued to respond, and the tracert was still able to reach the destination successfully. So, mostly the reasons could be the following:

- ➔ The ICMP responses for diagnostics using tracert may be considered of lower priority than the actual traffic data on that router, and hence may be dropped.
- ➔ A router may forward the packet on, but the ICMP response it generates may take a different return path that doesn't make it back to us.
- ➔ Some routers may be configured to ignore TTL-expired packets entirely to avoid exposing their internal topology.

tracert and captured packets in Wireshark for Windows:

```
Tracing route to www.dev.to [172.67.69.135]
over a maximum of 30 hops:
```

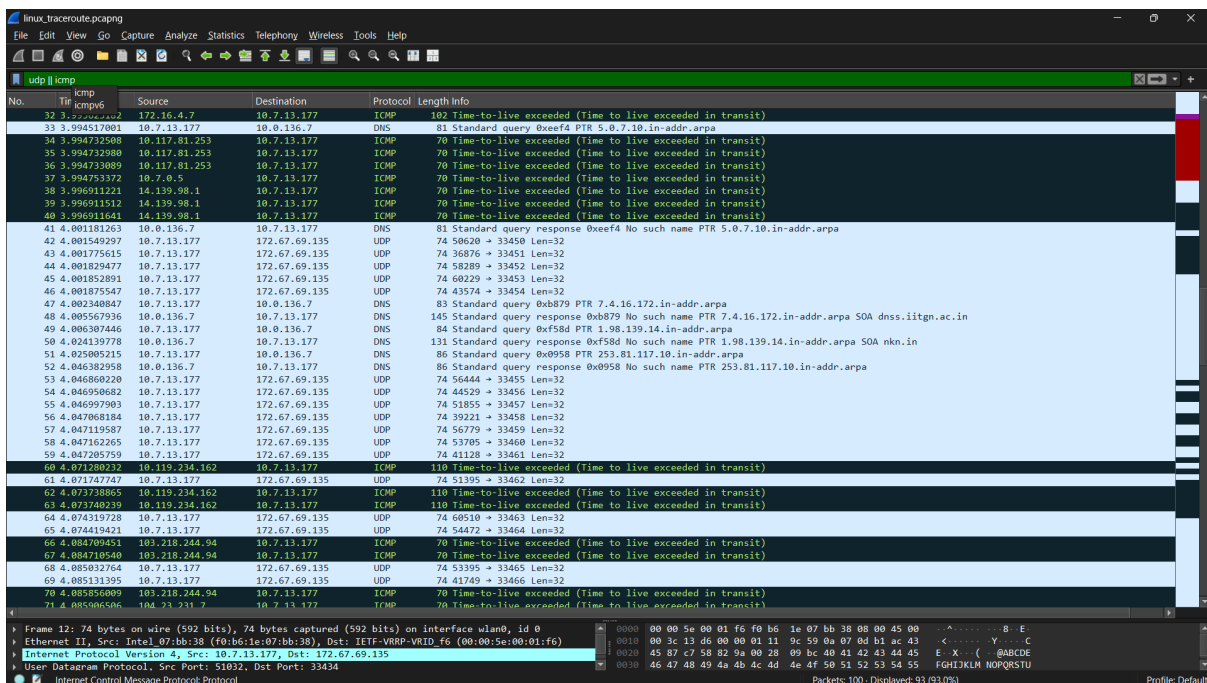
```
  0  0 ms  0 ms  0 ms  10.7.0.5
  1  10 ms  2 ms  2 ms  172.16.4.7
  2  15 ms  4 ms  3 ms  14.139.98.1
  3  10 ms  2 ms  2 ms  10.117.81.253
  4  *      *      *      Request timed out.
  5  *      *      *      Request timed out.
  6  *      *      *      Request timed out.
  7  21 ms  23 ms  22 ms  10.119.234.162
  8  41 ms  50 ms  47 ms  103.218.244.94
  9  52 ms  34 ms  33 ms  104.23.231.7
 10  37 ms  32 ms  32 ms  172.67.69.135
```

Trace complete.

```
366 33.017140 10.7.48.124 172.67.69.135 ICMP 106 Echo (ping) request id=0x0001, seq=19/4864, ttl=5 (no response found!)
378 36.720374 10.7.48.124 172.67.69.135 ICMP 106 Echo (ping) request id=0x0001, seq=20/5120, ttl=5 (no response found!)
384 40.719946 10.7.48.124 172.67.69.135 ICMP 106 Echo (ping) request id=0x0001, seq=21/5376, ttl=5 (no response found!)
391 44.726594 10.7.48.124 172.67.69.135 ICMP 106 Echo (ping) request id=0x0001, seq=22/5632, ttl=6 (no response found!)
460 48.719964 10.7.48.124 172.67.69.135 ICMP 106 Echo (ping) request id=0x0001, seq=23/5888, ttl=6 (no response found!)
508 52.712599 10.7.48.124 172.67.69.135 ICMP 106 Echo (ping) request id=0x0001, seq=24/6144, ttl=6 (no response found!)
594 56.723420 10.7.48.124 172.67.69.135 ICMP 106 Echo (ping) request id=0x0001, seq=25/6400, ttl=7 (no response found!)
612 60.710026 10.7.48.124 172.67.69.135 ICMP 106 Echo (ping) request id=0x0001, seq=26/6656, ttl=7 (no response found!)
625 64.723878 10.7.48.124 172.67.69.135 ICMP 106 Echo (ping) request id=0x0001, seq=27/6912, ttl=7 (no response found!)
631 68.712645 10.7.48.124 172.67.69.135 ICMP 106 Echo (ping) request id=0x0001, seq=28/7168, ttl=8 (no response found!)
632 68.734196 10.119.234.162 10.7.48.124 ICMP 110 Time-to-live exceeded (Time to live exceeded in transit)
```

traceroute and captured packets in Wireshark for Linux:

```
> traceroute www.dev.to
traceroute to www.dev.to (172.67.69.135), 30 hops max, 60 byte packets
 0 10.7.0.5 (10.7.0.5) 2.352 ms 2.180 ms 3.247 ms
 1 172.16.4.7 (172.16.4.7) 2.059 ms 2.000 ms 1.945 ms
 2 14.139.98.1 (14.139.98.1) 5.177 ms 5.046 ms 4.979 ms
 3 10.117.81.253 (10.117.81.253) 2.743 ms 2.690 ms 2.597 ms
 4 * * *
 5 * * *
 6 * * *
 7 * * *
 8 10.119.234.162 (10.119.234.162) 24.501 ms 26.807 ms 26.757 ms
 9 103.218.244.94 (103.218.244.94) 37.680 ms 37.607 ms 38.709 ms
10 104.23.231.7 (104.23.231.7) 38.716 ms 104.23.231.5 (104.23.231.5) 34.810 ms 104.23.231.7 (104.23.231.7) 34.238 ms
11 172.67.69.135 (172.67.69.135) 36.295 ms 33.222 ms 36.856 ms
```



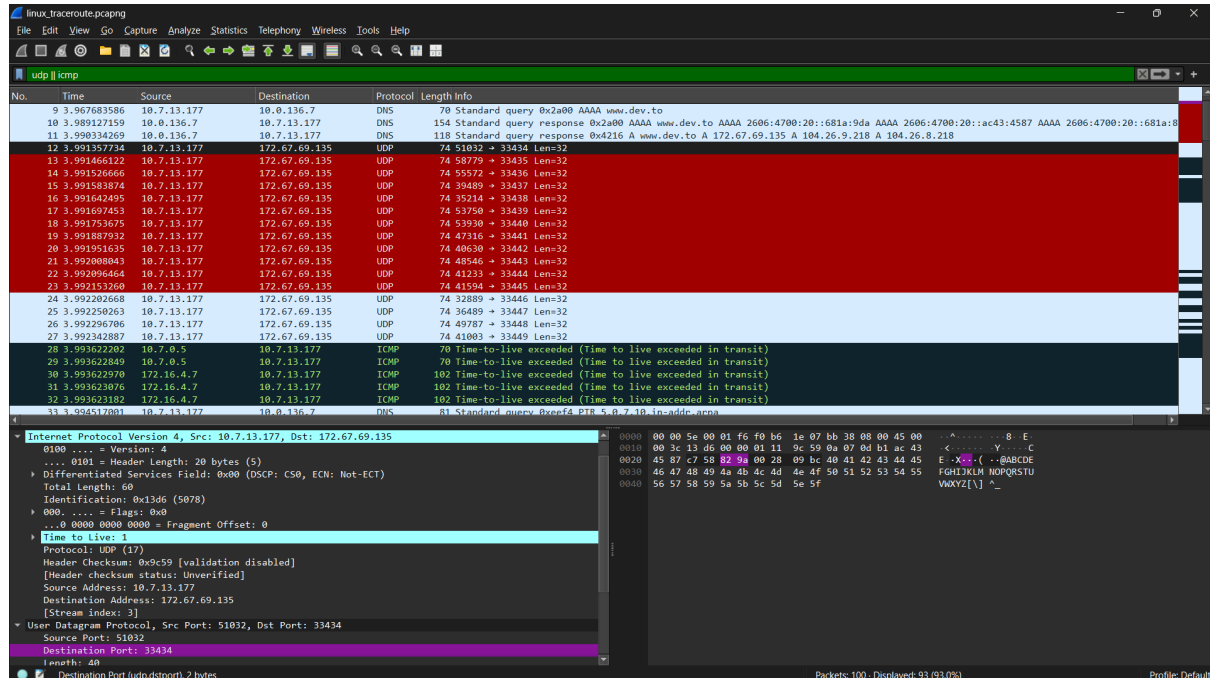
Here we can see that the packets 34, 35, 36 are received from IP 10.117.81.253, which is at hop 4, giving the ICMP TTL exceeded message, and then there is a large gap before it receives the same from the IP of hop 8.



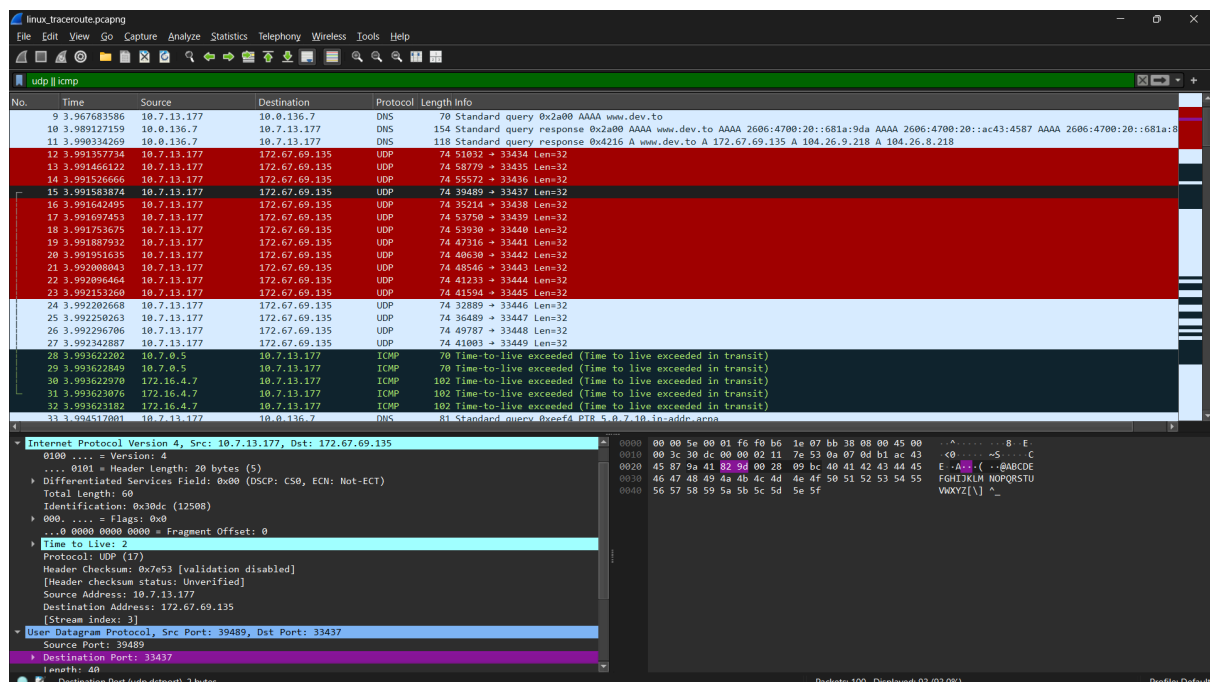
### 3. In Linux traceroute, which field in the probe packets changes between successive probes sent to the destination?

Ans. In Linux traceroute, two fields change:

- The TTL (Time-To-Live) field → increments for each hop (1, 2, 3, ...).
- The UDP destination port number → changes for each probe, starting at 33434 and incrementing (33435, 33436, ...).



In the above image, we can see that the TTL is 1 and the destination port is 33434 (this is for packet 12, which is the first UDP packet sent of hop 1). And in the below image, we can see that the TTL is 2 and the destination port is 33435 (which is for packet 15, corresponding to the first packet sent of hop 2).



**4. At the final hop, how is the response different compared to the intermediate hop?**

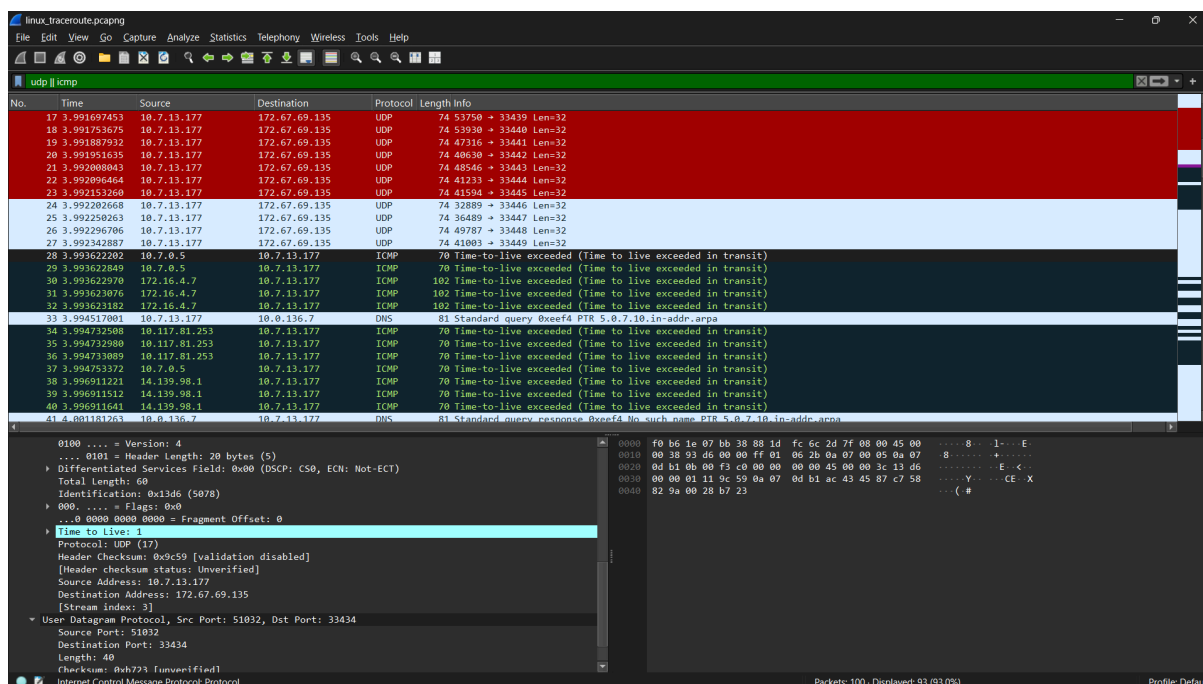
Ans. In Windows, at the final hop, each ICMP echo request of the final hop receives an echo reply from the destination IP.

869	85.607921	10.7.48.124	172.67.69.135	ICMP	106	Echo (ping) request	id=0x0001, seq=37/9472, ttl=11 (reply in 870)
870	85.645357	172.67.69.135	10.7.48.124	ICMP	106	Echo (ping) reply	id=0x0001, seq=37/9472, ttl=53 (request in 869)
871	85.646247	10.7.48.124	172.67.69.135	ICMP	106	Echo (ping) request	id=0x0001, seq=38/9728, ttl=11 (reply in 872)
872	85.678630	172.67.69.135	10.7.48.124	ICMP	106	Echo (ping) reply	id=0x0001, seq=38/9728, ttl=53 (request in 871)
873	85.679834	10.7.48.124	172.67.69.135	ICMP	106	Echo (ping) request	id=0x0001, seq=39/9984, ttl=11 (reply in 874)
874	85.711903	172.67.69.135	10.7.48.124	ICMP	106	Echo (ping) reply	id=0x0001, seq=39/9984, ttl=53 (request in 873)

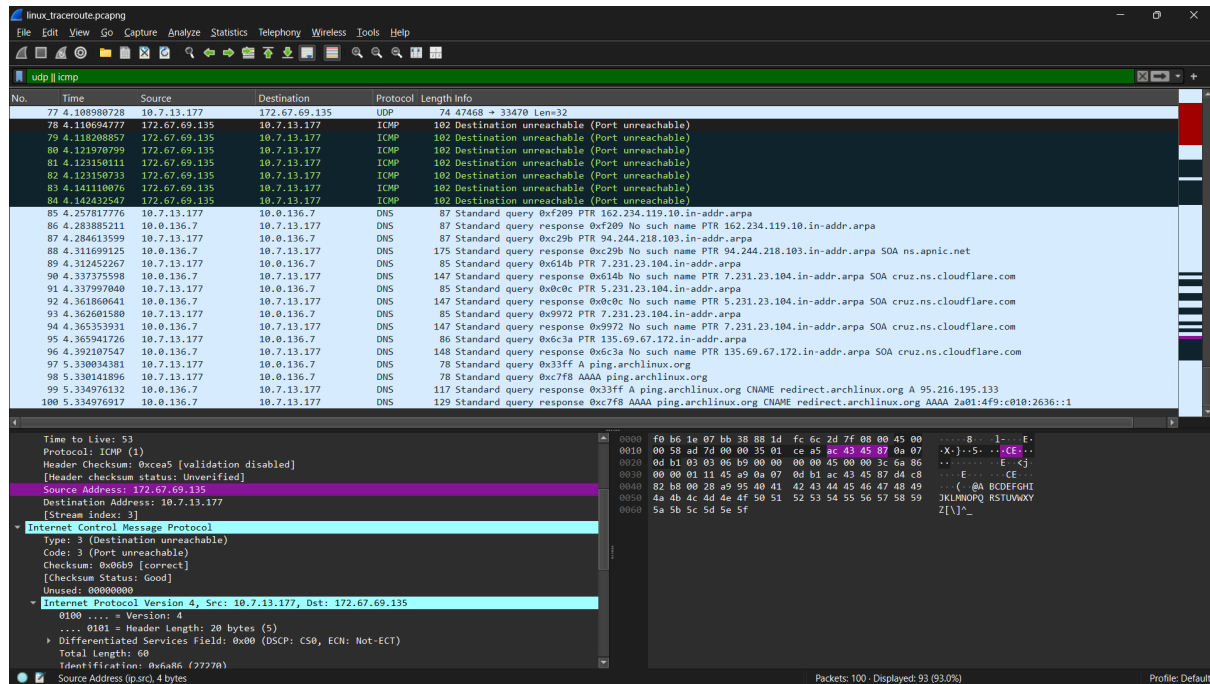
Whereas at the intermediate hops, each ICMP echo request was followed by an ICMP error message saying “Time-to-live exceeded”, as the packets didn’t reach the destination within the fixed TTL.

126	16.303820	10.7.48.124	172.67.69.135	ICMP	106 Echo (ping) request id=0x0001, seq=10/2560, ttl=2 (no response found!)
127	16.314366	172.16.4.7	10.7.48.124	ICMP	134 Time-to-live exceeded (Time to live exceeded in transit)
128	16.316690	10.7.48.124	172.67.69.135	ICMP	106 Echo (ping) request id=0x0001, seq=11/2816, ttl=2 (no response found!)
129	16.318724	172.16.4.7	10.7.48.124	ICMP	134 Time-to-live exceeded (Time to live exceeded in transit)
130	16.319976	10.7.48.124	172.67.69.135	ICMP	106 Echo (ping) request id=0x0001, seq=12/3072, ttl=2 (no response found!)
131	16.321826	172.16.4.7	10.7.48.124	ICMP	134 Time-to-live exceeded (Time to live exceeded in transit)
145	21.889113	10.7.48.124	172.67.69.135	ICMP	106 Echo (ping) request id=0x0001, seq=13/3328, ttl=3 (no response found!)
146	21.904199	14.139.98.1	10.7.48.124	ICMP	70 Time-to-live exceeded (Time to live exceeded in transit)
147	21.906945	10.7.48.124	172.67.69.135	ICMP	106 Echo (ping) request id=0x0001, seq=14/3584, ttl=3 (no response found!)
148	21.911332	14.139.98.1	10.7.48.124	ICMP	70 Time-to-live exceeded (Time to live exceeded in transit)
149	21.912971	10.7.48.124	172.67.69.135	ICMP	106 Echo (ping) request id=0x0001, seq=15/3840, ttl=3 (no response found!)
150	21.916788	14.139.98.1	10.7.48.124	ICMP	70 Time-to-live exceeded (Time to live exceeded in transit)

Whereas in Linux, at intermediate hops, routers send back ICMP TTL exceeded messages when the TTL reaches zero. We can see that in the image below that packet 28 is the ICMP TTL exceeded message received for our first UDP packet sent on hop 1, which had a TTL of 1 and destination port 33434.



And at the final hop (destination), since the UDP probe reaches the host but no service is listening on the high port, the host replies with an ICMP Port Unreachable error message (Type 3, Code 3). We can see that in the image below, the packet 74 comes straight from our target IP.



**5. Suppose a firewall blocks UDP traffic but allows ICMP — how would this affect the results of Linux traceroute vs. Windows tracert?**

Ans.

- Linux traceroute (UDP-based): The probes would never reach the destination. No ICMP Port Unreachable responses would come back, so traceroute would fail or hang.
- Windows tracert (ICMP-based): It would still work correctly, since ICMP Echo Requests and Replies are allowed through the firewall.

Thus, Windows tracert is more resilient in this scenario.